

Exploratory Cluster Analysis for Josie Schafer

Michael Flynn, Prior Analytics, LLC.

Cleaning Data

This first code block loads the data and performs any necessary cleaning, rescaling, etc.

First, there don't appear to be any missing values in any rows.

Second, for now I'm focusing primarily on broader demographic and institutional indicators for now, but also some more targeted variables that would likely help to explain disparate economic outcomes. For example, high-research universities granting PhDs or a higher number of community hospital beds. There are other variables that we could include that might be useful for some purposes (e.g. Medicare recipients by region) but I expect that these will be closely tracking other age-related demographic variables.

Third, I'm rescaling the variables by dividing the observed value by the largest value of X as follows:

$$\frac{X_i}{\max X}$$

This puts all observed values on a 0 – 1 scale.

My understanding of clustering techniques is that when they calculate the distance between units, they will treat the scale of the variables equivalently. The idea here is to scale all of the cluster inputs so they are all on a 0 – 1 scale, thereby treating all of them equivalently. That way variables with large values and large ranges don't dominate the clustering procedure.

That said, if there's reason to want to weight input variables differently for clustering we can explore that with more time.

```
# Read in data and select relevant variables for clusters.  
# Focus is on demographic and anchor institution variables.  
  
# Read in raw data file  
data <- readxl::read_xlsx(here("data/anchor regions analysis.xlsx"))
```

```

# List of variables to include in clustering
varlist <- c("totpop_19", # Total pop
            "popchange", # pop change
            "medage",    # Median age
            "labfor",    # percent population in labor force
            "pov",       # percent population living in poverty
            "poc",       # people of color as percent of pop
            "highed",    # Percent population with at least bachelor's
            "forborn",   # Percent population foreign born
            "net_mig",   # Net domestic migration
            "highered_emp_qcew",
            "highered_estab_qcew",
            "hospital_emp_qcew",
            "hospital_estab_qcew",
            "inst_ipeds_enrollment_all",
            "inst_ipeds_doctoralunihighrese",
            "inst_ipeds_pellawards",
            "inst_hosp_ahacommunityhospitals",
            "inst_hosp_ahabeds",
            "inst_hosp_nihresearchfunding")

# Rescale the variables from 0-1
data.clean <- data |>
  mutate(across(all_of(varlist), # Variables to scale
               ~.x/max(.x),      # Scale relative to
               .names = "{col}_max")) |> # Add "max" suffix
  dplyr::select(MSA, ends_with("_max")) |> # select chosen vari
  column_to_rownames("MSA")

```

Clustering Methods

Here I start with agglomerative/hierarchical clustering methods. The goal as I understand it is to find a happy medium number of groups that illustrates the variability across regions and anchor institutions while still being tractable for analyses.

The priority here is to construct clusters on the basis of 1) anchor institution characteristics, and 2) demographic characteristics of the surrounding region. For now I'll combine these into a single cluster, but we may want to think about constructing two clusters, one on the basis of demographic traits and the other on the basis of anchor institution traits. This would help parse out effects later if the client is interested in using these as predictors in subsequent

regression analyses.

I'm going to create a few different clusters and we can compare the characteristics and performance of each, and then choose which one the client likes best.

I chose the “complete” method for the `hclust()` function because it generates a better distribution of clusters than the other methods. For example, others tend to produce either very flat distributions, in which case you may just as well use dummy variables for each MSA or city, or they produce oddly concentrated clusters with 80-90% of observations falling into cluster group 2.

```
distance <- dist(data.clean) # calculate Euclidian distance between obs

hc.tree <- hclust(distance, method = "complete") # Create cluster groupings based on distance

# List of distances to use in generating clusters
cluster.size.list <- list("5" = 5,
                          "10" = 10,
                          "15" = 15,
                          "20" = 20,
                          "25" = 25,
                          "30" = 30,
                          "35" = 35,
                          "40" = 40)

cluster.ids <- map(
  .x = seq_along(cluster.size.list),
  .f = ~ cutree(hc.tree, k = cluster.size.list[[.x]])
) |>
  bind_cols()
```

New names:

```
* `` -> `...1`
* `` -> `...2`
* `` -> `...3`
* `` -> `...4`
* `` -> `...5`
* `` -> `...6`
* `` -> `...7`
* `` -> `...8`
```

```
names(cluster.ids) <- c("cluster_5", "cluster_10", "cluster_15", "cluster_20", "cluster_25")

data.out <- data.clean |>
  bind_cols(cluster.ids)
```

Choosing the optimal number of clusters

Figure 1 shows the distribution of the observations depending on the number of clusters chosen. In general, 25–35 clusters seems like a nice balance between parsimony and too much detail. Smaller numbers of clusters, like 5 or 10, group too many areas together (see the spike at group #1). In general we see there are regularly spikes like these, but we start to get more variability as we move towards the 25–30 range.

```
cluster.list <- list("cluster_5", "cluster_10", "cluster_15", "cluster_20", "cluster_25",

plot.out <- data.out |>
  dplyr::select(starts_with("cluster")) |>
  map2(
    .y = cluster.list,
    .f = ~{ggplot(data.out, aes(x = .)) +
      geom_histogram(bins = 40) +
      labs(title = .y)}
  )

patchwork::wrap_plots(plot.out)
```




Figure 1: Histograms showing the distribution of clusters depending on the number of clusters chosen.

```
table.data <- data |>
  bind_cols(cluster.ids) |>
  dplyr::select(varlist, cluster_30) |>
  group_by(cluster_30) |>
  dplyr::summarise(across(everything(),
                           mean))
```

Warning: Using an external vector in selections was deprecated in tidysselect 1.1.0.
i Please use `all_of()` or `any_of()` instead.

Was:

```
data %>% select(varlist)
```

Now:

```
data %>% select(all_of(varlist))
```

See <<https://tidysselect.r-lib.org/reference/faq-external-vector.html>>.

```
table.out <- table.data |>
  kbl(longtable = TRUE) |>
  kable_styling(font_size = 8) |>
```

```
scroll_box(height = "600px", width = "800px")
```

```
table.out
```

cluster_30	totpop_19	popchange	medage	labfor	pov	poc	highed	forborn	net_mig	higherec
1	491022.2	4.987475	38.27273	64.10909	11.21818	32.31163	33.54545	21.427273	-1.448091e+03	
2	994230.7	7.819001	37.26667	63.83333	14.40000	38.49770	31.41111	6.000000	-4.233333e+01	
3	214126.1	4.060407	39.85067	58.44000	15.88533	17.23172	22.28400	4.241333	1.025733e+02	
4	2319054.6	19.071504	36.69000	67.13000	11.48000	26.89717	37.07000	14.100000	1.279350e+04	
5	246575.0	7.490234	38.01528	64.73333	11.86806	13.37229	26.31111	6.188889	-6.861111e+00	
6	577068.9	8.442607	37.72586	65.01379	12.75517	17.88891	35.38966	7.372414	-2.206897e+01	
7	253319.8	10.775605	32.23077	61.05385	19.19231	20.70978	41.03077	9.261539	-6.108462e+02	
8	299882.5	9.290796	35.95000	59.55769	19.82692	43.58506	25.08077	3.815385	-2.716538e+02	
9	447905.1	8.002275	32.08667	58.36000	22.06000	23.57170	17.84000	22.613333	-5.049333e+02	
10	575400.0	2.552712	39.55000	66.70000	8.80000	71.78968	31.10000	19.200000	-3.494500e+03	
11	342088.6	10.687640	50.82727	47.78182	14.18182	15.21998	22.47273	8.936364	4.332455e+03	
12	125044.0	0.000000	67.40000	22.50000	8.20000	10.99933	31.80000	5.300000	2.779000e+03	
13	637543.3	18.178695	50.60000	52.03333	11.53333	12.83993	32.46667	18.200000	1.163967e+04	
14	328124.3	9.772405	39.70000	66.56667	10.50000	11.81458	48.58333	8.033333	-2.433333e+01	
15	2197690.2	6.524644	37.91000	67.07000	11.60000	22.13352	35.52000	7.210000	-2.119200e+03	
16	404417.0	227.829477	45.50000	57.40000	12.60000	23.64070	27.30000	7.100000	2.044000e+03	
17	386114.5	72.536067	40.30000	58.55000	17.30000	26.97233	23.12500	3.800000	1.480000e+03	
18	2821709.5	31.846792	39.80000	63.45000	12.60000	27.68355	32.75000	12.000000	1.957450e+04	
19	3461420.8	6.809083	37.22000	63.52000	12.88000	34.85319	33.10000	16.740000	-7.903600e+03	
20	6090660.0	11.166374	41.00000	63.00000	14.60000	29.79406	32.30000	40.700000	4.355000e+03	
21	4761603.0	16.685736	36.70000	62.80000	13.70000	22.16783	31.50000	14.300000	4.338300e+04	
22	3344589.0	10.782462	38.05000	67.55000	8.25000	52.68919	50.60000	34.650000	-2.523000e+03	
23	6373281.0	17.481792	35.35000	66.95000	12.90000	40.77722	35.70000	18.600000	1.343550e+04	
24	6196585.0	14.397978	37.00000	71.60000	7.80000	46.53421	50.90000	22.800000	-1.732200e+04	
25	7320663.0	18.952678	34.80000	68.80000	11.70000	31.66168	35.20000	18.700000	3.216400e+04	
26	4832346.0	7.642613	38.70000	69.20000	9.30000	23.97910	48.10000	18.900000	-1.439600e+04	
27	6079130.0	2.833259	38.80000	65.30000	12.40000	33.39762	37.90000	11.000000	-1.392500e+04	
28	9508605.0	1.320708	37.50000	66.80000	11.80000	34.26053	38.00000	17.700000	-4.973800e+04	
29	13249614.0	4.132679	36.80000	64.90000	13.90000	46.41575	34.50000	33.100000	-3.615400e+04	
30	19294236.0	3.173788	38.60000	64.70000	12.80000	42.45710	40.40000	29.500000	-1.382570e+05	